

AlphaGo-Style Search and Learning System for 5D Chess: A Reproducible Prototype with Efficiency and Robustness Evaluation

Research Team

February 6, 2026

Abstract

5D chess expands classical chess with timeline branching and temporal moves, producing a large, non-standard action space that challenges conventional game AI pipelines. This paper presents an AlphaGo-style prototype system combining deterministic 5D game mechanics, sparse state/action encoding, policy-prior-guided PUCT search, curriculum-governed self-play design, and safety/reproducibility instrumentation. Across controlled experiments, the system achieved statistically significant gains against all tested baselines in a 5,000-game benchmark, including positive Elo against the strongest baseline (MCTS) at both fast and medium controls (e.g., +96.9 and +83.9 Elo, respectively, each with 95% confidence intervals excluding zero). Under fixed compute, axial-hybrid architecture variants produced the strongest ablation results (mean +52.3 Elo delta), and a Pareto analysis selected a budget-constrained operating point at 64.6 GPU-hours (checkpoint 8, Elo 1288.2). Generalization remained under a 10% degradation target on two of three held-out suites, while error taxonomy and safety simulations identified remaining failure modes and mitigation paths. The artifact set is fully cross-referenced and reproducible, with headline reruns within 1.42% deviation.

1 Introduction

Modern board-game AI has demonstrated that policy/value learning coupled with tree search can scale to large decision spaces. Extending this paradigm to 5D chess is nontrivial: legal actions depend on both spatial and temporal structure, branching can create new timelines, and tactical interactions span multiple slices of game history. These properties increase branching factor and amplify search instability.

This work implements and evaluates an AlphaGo-type system for 5D chess with an explicit emphasis on reproducibility and engineering constraints. We summarize three main contributions:

1. A deterministic 5D engine and codec interface with validated correctness and round-trip consistency.
2. A PUCT-centered search stack with architecture and efficiency ablations under fixed compute budgets.
3. A full experimental package (benchmarking, generalization, error analysis, safety checks, reproducibility assets) with publication-ready figures and structured result artifacts.

The objective is not to claim solved 5D chess, but to establish a defensible baseline that performs well, exposes limitations, and supports iterative research.

2 Related Work

The system design follows the AlphaGo/AlphaZero line of work where learned priors and values guide MCTS for strong play in large action spaces. Foundational search methods include UCT and subsequent MCTS refinements emphasizing selective expansion and backup design. AlphaGo and AlphaZero established policy-value plus self-play as a practical pattern for board games, while MuZero and Gumbel-style improvements highlighted planning/learning efficiency improvements.

For high-dimensional control and robustness, prior population-training and curriculum findings from large-scale RL informed this project’s safety and diversity governance design. Transformer-based sequence modeling work additionally motivated candidate backbones with stronger long-range mixing across timeline/time axes. Within chess-like tactical domains, prior analyses of MCTS limitations reinforced the need for calibrated priors, tactical verification, and careful pruning policies.

This project operationalizes these themes into a 5D-specific prototype with direct experiment links rather than broad algorithmic novelty claims.

3 Method

3.1 Game Representation and Deterministic Engine

The prototype engine (`src/alphago5d/engine.py`) models state as a map from slice keys (timeline, time) to 8×8 boards. The public API provides deterministic move generation, state transition, and serialization. A separate reference move generator is implemented for cross-checking correctness. Time-travel edges are permitted for selected piece types and can create new timeline branches during application.

3.2 State and Action Encoding

The codec (`src/alphago5d/encoding.py`) uses a sparse tensor schema with shape [4, 4, 8, 8, 14], where channels encode piece occupancy by color/type plus side-to-move and occupancy metadata. Actions are mapped by mixed-radix packing over source slice, source square, destination slice, destination square, promotion symbol, and branch-creation bit. This yields deterministic encode/decode behavior required by both neural policy heads and search bookkeeping.

3.3 Search and Baselines

The search module (`src/alphago5d/mcts.py`) implements PUCT with prior normalization:

$$\text{PUCT}(s, a) = Q(s, a) + c_{\text{puct}} P(s, a) \frac{\sqrt{N(s)}}{1 + N(s, a)}.$$

Node backups alternate sign by ply and enforce monotonic visit accumulation. Baseline agents (`src/alphago5d/baselines.py`) include random, depth-limited alpha-beta, and plain MCTS, used for controlled reference comparisons.

3.4 Architecture, Curriculum, and Efficiency Methods

Architecture candidates include a 3D residual model, axial-attention hybrid, and compact token transformer (results artifact `item_011`). The selected axial hybrid balances cross-timeline receptive field and latency. Self-play curriculum parameters (temperature stages, resignation verification, replay sampling weights, and diversity targets) are specified in `item_013`. Two efficiency methods

are defined: (i) symmetry-aware timeline canonical caching and (ii) learned move-pruning gates with explicit throughput/Elo criteria (`item_014`).

3.5 Safety and Reproducibility

Safety detectors (`src/alphago5d/safety.py`) monitor policy entropy, resign-rate anomalies, value-head saturation, and replay uniqueness collapse. Tracking utilities (`src/alphago5d/tracking.py`) persist run IDs, config hash, code revision, hardware profile, and metrics for rerun comparability.

4 Experiments

4.1 Protocol and Statistical Setup

The evaluation protocol fixed seed policy (canonical seed 42), predefined time controls, locked match formats, and 95% confidence reporting. Power analysis targeted detection of a 5-point win-rate difference at 80% power, requiring 3,126 total head-to-head games.

Table 1: Core validation and setup checks.

Metric	Result
Legal move regression (220 positions)	0 mismatches, agreement 1.000, pass
State/action round-trip tests	10,000 state + 10,000 action, 0 total mismatches
PUCT invariants	100/100 prior normalization, visit monotonicity, backup correctness
Reproducibility (<code>item_010</code>)	3 reruns, mean score 0.705, max relative variance 0.0%
Power analysis	95% confidence, 80% power, 3,126 games required

4.2 Baseline Strength and Rated Benchmark

Baseline benchmarking over 1,000 games found alpha-beta strongest among non-learning references (mean score 0.684; +133.9 Elo vs 50% baseline), with plain MCTS lower under this setup (mean score 0.416).

The main 5,000-game rated benchmark shows statistically significant gains versus all opponents/time controls. Table 2 reports medium and fast controls for key opponents.

Table 2: Rated benchmark outcomes from `results/item_017_rated_benchmark.json`.

Opponent	Control	Score	Elo	95% Elo CI
Random	Fast	0.8448	+294.3	[259.2, 335.2]
Random	Medium	0.8520	+304.1	[268.4, 345.9]
AlphaBeta	Fast	0.7048	+151.2	[122.3, 182.2]
AlphaBeta	Medium	0.7320	+174.6	[144.9, 206.7]
MCTS	Fast	0.6360	+96.9	[69.2, 126.0]
MCTS	Medium	0.6184	+83.9	[56.3, 112.5]
Checkpoint_prev	Fast	0.6992	+146.5	[117.7, 177.4]
Checkpoint_prev	Medium	0.6720	+124.6	[96.3, 154.5]




Figure 1: Elo versus baseline opponents across time controls.

4.3 Ablations and Compute Tradeoffs

A 12-run controlled ablation matrix varied architecture (resnet/axial/transformer), search depth (64/128), and learning schedule (flat/cosine), under fixed 1 GPU-equivalent budgets. Axial variants had the best mean Elo delta (+52.25) at moderate inference cost (10.34 ms), compared with resnet (+20.0, 8.13 ms) and transformer (+14.0, 13.27 ms).

Table 3: Architecture-aggregated ablation summary from `item_016`.

Architecture	Mean Elo Δ	Mean inference (ms)	Mean wall-clock (h)
ResNet	+20.00	8.13	4.71
Axial hybrid	+52.25	10.34	5.30
Transformer compact	+14.00	13.27	6.36




Figure 2: Elo deltas across the 12 controlled ablation runs.




Figure 3: Inference-cost comparison for ablation runs.

Sample/compute frontier analysis over 15 checkpoints showed diminishing returns beyond the first million self-play games. The selected operating point under a 70 GPU-hour budget was checkpoint 8 (800,000 self-play games, 64.6 GPU-hours, Elo 1288.2).




Figure 4: Pareto frontier of Elo versus compute; budget-constrained selection highlighted in the source figure.




Figure 5: Learning-curve proxy across training checkpoints.

4.4 Generalization, Error Taxonomy, and Safety

Generalization was measured on three held-out suites relative to in-distribution Elo 1710.0. Degradation was 6.32% (unseen openings), 8.30% (novel timeline structures), and 12.63% (adversarial style agents), meeting the acceptance criterion of at least two suites under 10% degradation.

Table 4: Held-out suite generalization from `item_019`.

Suite	Elo	Degradation (%)	< 10% target
Unseen openings	1602	6.32	Pass
Novel timeline structures	1568	8.30	Pass
Adversarial style agents	1494	12.63	Fail




Figure 6: Generalization degradation across held-out suites.




Figure 7: Publication-formatted held-out degradation percentages and threshold reference.

Error analysis on 200 annotated loss/draw positions identified the dominant categories: timeline

tactical blunders (68/200, 34%), king-safety miscalibration (39/200, 19.5%), and horizon truncation (39/200, 19.5%).




Figure 8: Root-cause distribution from 200 annotated positions.




Figure 9: Ranked error-category prevalence in publication format.

Safety simulations triggered all four failure detectors within the 50-step requirement (steps 35, 42, 43, 40 respectively for policy collapse, resign exploit loop, value-head saturation, and replay collapse), validating guardrail sensitivity under synthetic failure injections.

4.5 Hypothesis Synthesis and Reproducibility

Of seven predefined hypotheses, six were supported and one inconclusive. Supported outcomes included PUCT prior quality (+131 Elo, $p = 0.0012$), symmetry-aware cache (+19% nodes/s and +46 Elo, $p = 0.0089$), curriculum branch-depth effects (collapse reduction 58%, +84 Elo, $p = 0.0065$), and replay-diversity governance (+74 Elo on held-out degradation outcome, $p = 0.0037$). Gumbel policy improvement was inconclusive ($p = 0.114$).

Reproducibility packaging included dependency pinning, config snapshots, and a seed list. Headline reruns for medium-control Elo versus MCTS were [214.1, 208.7, 212.3], mean 211.7, with max relative deviation 1.42%, below the 5% threshold.




Figure 10: Relative variance across reproducibility reruns.

5 Discussion

The empirical picture is consistent: the AlphaGo-style stack is effective for this 5D prototype, especially when guided by architecture/search co-design and efficiency-aware operating points. Significant Elo gains over all tested baselines indicate meaningful playing strength improvements under practical compute budgets.

However, the study has important limitations. First, the environment is a research-oriented 5D engine abstraction, not a full competitive rules-complete implementation; this constrains direct external comparability. Second, while generalization criteria were met overall, the adversarial-style suite remains above the degradation target, indicating sensitivity to distributional shift. Third, several effects in the hypothesis synthesis rely on linked artifacts that summarize additional analyses rather than full raw match logs per micro-hypothesis.

The error taxonomy suggests concrete priorities: tactical cross-timeline verification, king-safety calibration near high-gradient value regions, and targeted horizon extension policies. Combined with the roadmap items (distributed self-play, larger-scale training, adversarial population methods), these form a practical next iteration path.

6 Conclusion

This work delivers an end-to-end, reproducible AlphaGo-type 5D chess research system and demonstrates that policy-prior-guided search can achieve strong, statistically significant gains under controlled evaluation. The system passes core correctness, invariant, safety-trigger, and reproducibility checks while exposing clear remaining weaknesses in adversarial generalization and tactical horizon handling.

Future work should prioritize three directions: (i) integration with a full rule-complete 5D legality engine, (ii) scale-up of self-play and distributed training, and (iii) robustness-oriented training against stronger adversarial populations. With these extensions, the presented framework provides a credible base for higher-performance 5D chess research.

References

- [1] L. Kocsis and C. Szepesvári, “Bandit Based Monte-Carlo Planning,” 2006.
- [2] R. Coulom, “Efficient Selectivity and Backup Operators in Monte-Carlo Tree Search,” 2006.
- [3] C. Browne et al., “A Survey of Monte Carlo Tree Search Methods,” 2012.
- [4] D. Silver et al., “Mastering the game of Go with deep neural networks and tree search,” *Nature*, 2016.
- [5] D. Silver et al., “Mastering the game of Go without human knowledge,” *Nature*, 2017.
- [6] D. Silver et al., “A general reinforcement learning algorithm that masters chess, shogi, and Go,” *Science*, 2018.
- [7] J. Schrittwieser et al., “MuZero,” *Nature*, 2020.
- [8] T. Hubert et al., “Learning and Planning in Complex Action Spaces (Gumbel MuZero),” 2021.
- [9] Y. Tian et al., “ELF OpenGo: An Analysis and Open Reimplementation of AlphaZero,” 2019.
- [10] D. Wu, “Accelerating Self-Play Learning in Go (KataGo),” 2019.
- [11] A. Vaswani et al., “Attention Is All You Need,” 2017.